



# Car Sales Forecasting


## Leveraging Machine Learning to Optimize Inventory and Maximize Profitability

**Date:** July 26, 2025

**Organization:** Data Science Institute - University of Toronto

**Presented by:** ML Team 3

# Business Problem & Objective



## The Challenge

A large dealership network struggled with accurately forecasting monthly sales in highly competitive markets.



## Our Solution

Our machine learning model predicts the **top 3 best-selling** vehicles each month using historical sales data and car attributes. This solution helps dealers to:

- Stock the right vehicles
- Reduce overstock inventory
- Increase overall profitability

# Dataset Overview



## Data Source

Car Sales Report dataset from Kaggle + 3 synthetic features  
~28,000 records over 24 months  
(Jan 2022 – Dec 2023)



## Key Features

- Car specifications
- Buyer information
- Dealer region
- Family Size (synthetic)
- Gas Mileage (synthetic)
- Crash Test Score (synthetic)



## Limitations

- **Limited data volume:** ~28K samples limit deep modeling.
- **Severe imbalance:** Only 7% labeled “best sellers”.
- **Compute constraints:** Limited to models with reasonable training times.



# Modeling Approach



## Data Preparation

Separated numeric and categorical features

Addressed class imbalance via weighting



## Model Selection

- Logistic Regression
- Random Forest
- XGBoost
- LightGBM



## Optimization

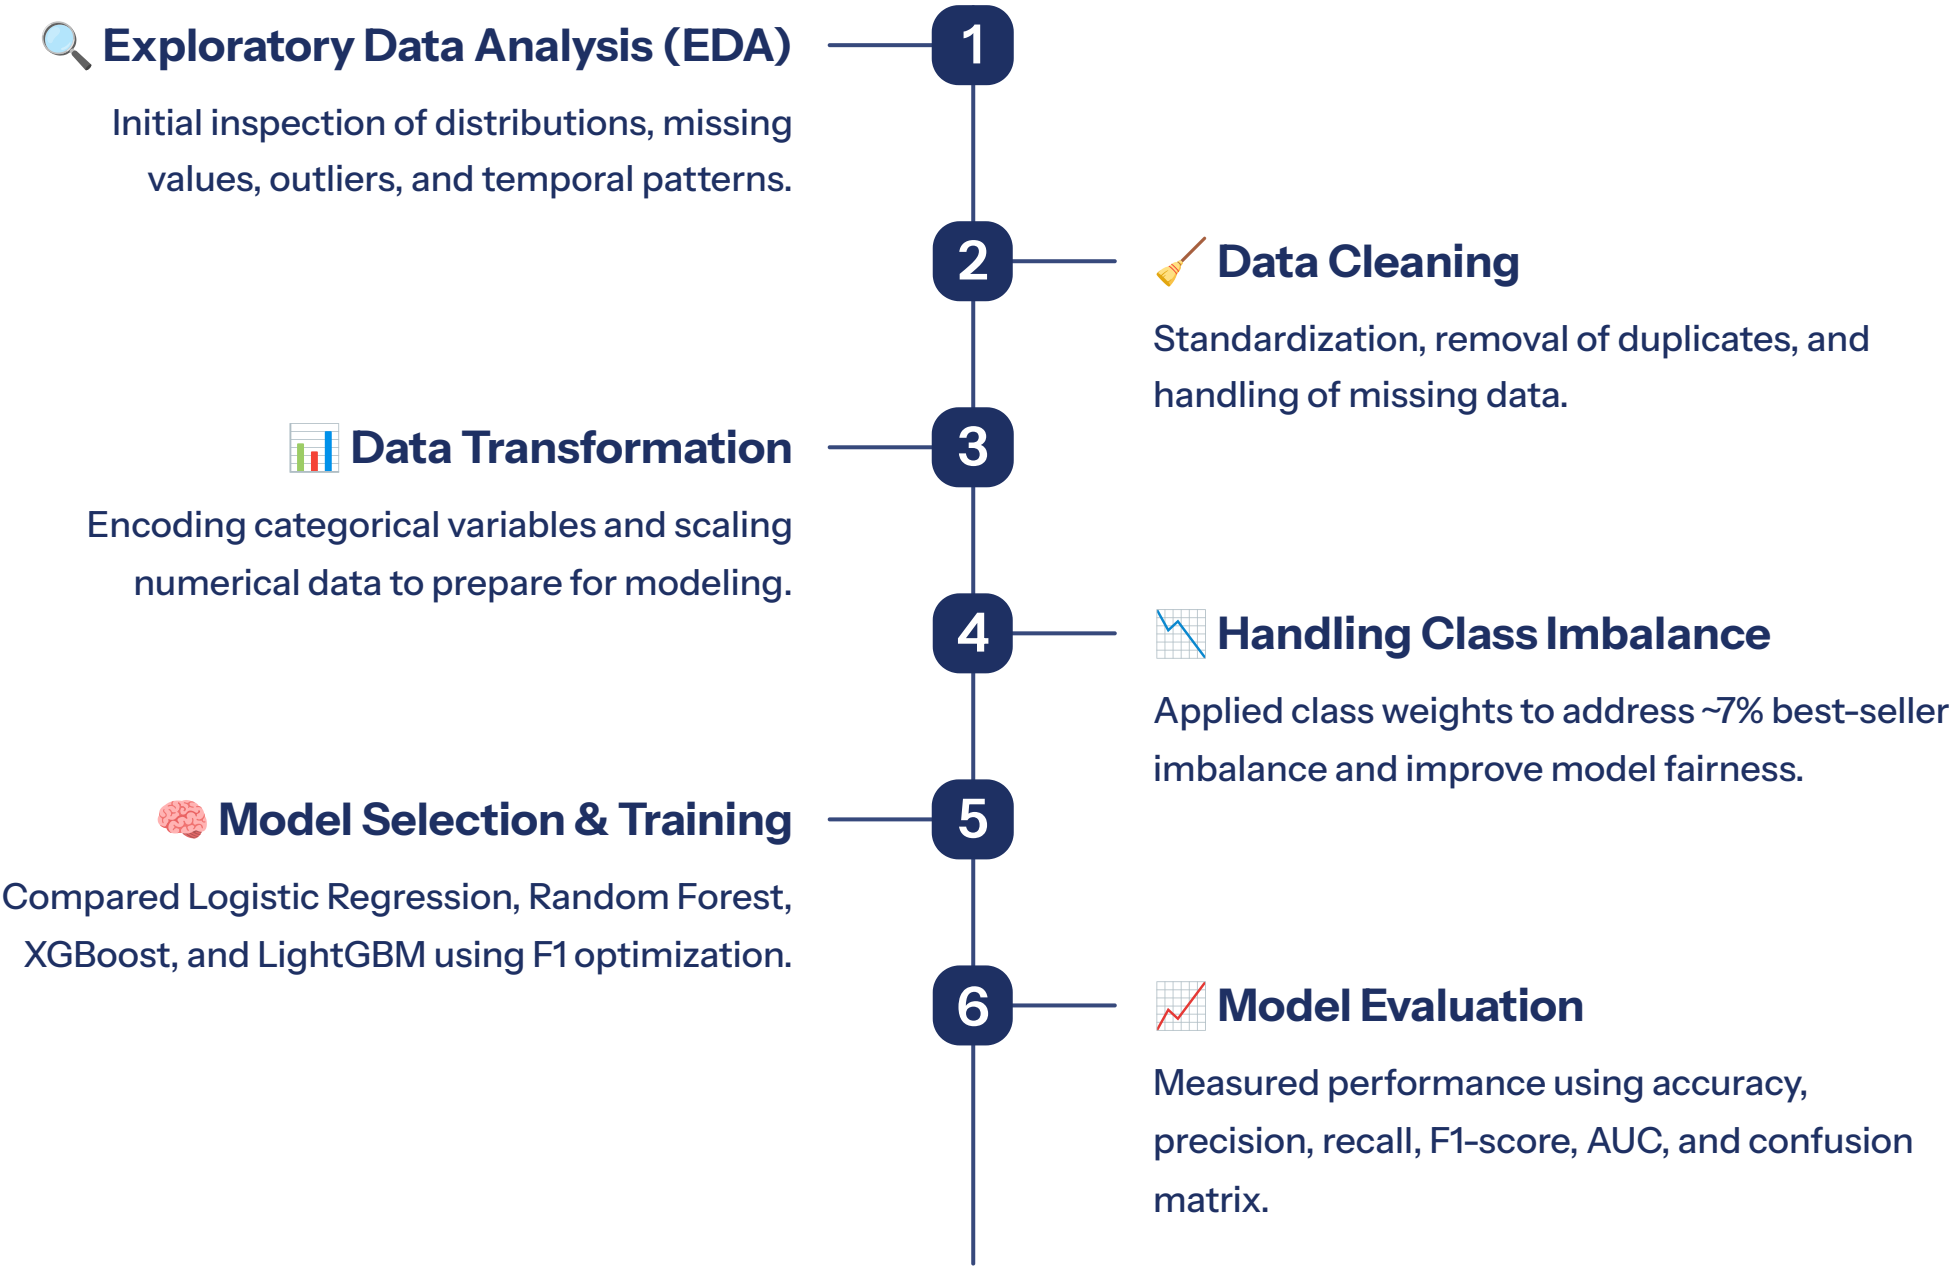
Grid search with F1-score as optimization metric

Test data = last 3 months to simulate real-world forecasting




# Analysis Methods and Tools


## Methodology Overview




## Tools Used




**Pandas**  
Utilized for powerful data manipulation and analysis, especially with structured data.



**NumPy**  
Leveraged for numerical computing, providing support for arrays and mathematical functions.



**Matplotlib & Seaborn**  
Employed for creating static, interactive, and animated visualizations to explore data patterns.



**Scikit-learn**  
The primary library for implementing various machine learning algorithms and predictive modeling.

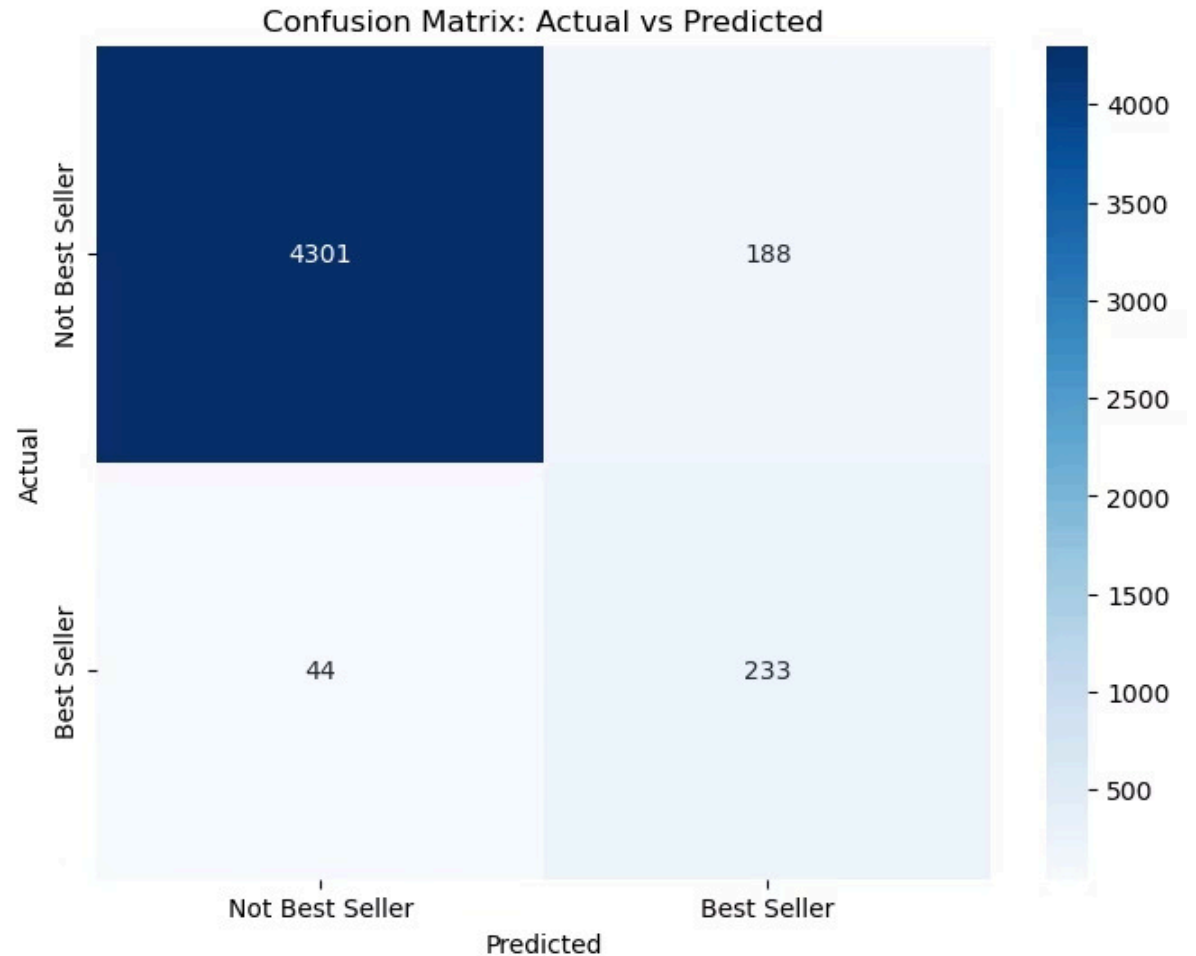
# Model Evaluation

## Performance Metrics


### Model: XGBoost (final selected)

- Accuracy: 0.9513
- Precision: 0.5534
- Recall: 0.8411
- F1 Score: 0.6676

Final model (XGBoost) selected based on balanced performance across precision and recall. High recall ensures that most best sellers are correctly identified, supporting inventory planning decisions.

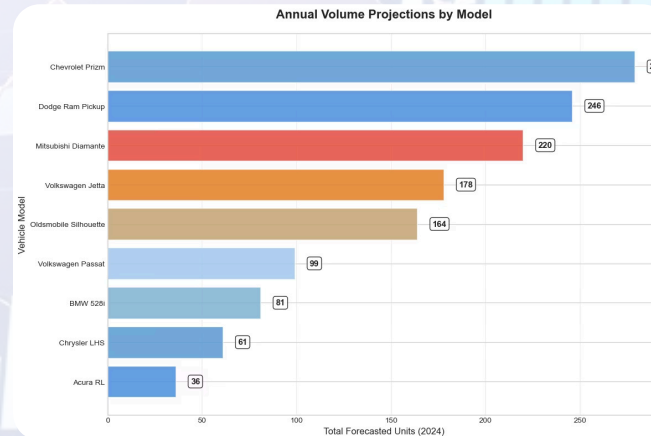


# Sales Forecast Highlights – Top Models and Inventory Planning




## Q4 2024 Forecast

Reveals seasonal trends and the most in-demand models to prepare for year-end sales peaks.



## Annual Projections

Identifies the highest-volume vehicles forecasted for 2024, enabling informed stocking decisions.



## Monthly Inventory Requirements

Supports monthly order planning by anticipating vehicle demand and minimizing overstock or stockouts.



# Business Impact

## Improved Inventory

Dealers can now anticipate best sellers monthly, optimizing stock levels.

## Higher Sales

Right cars at the right time leads to increased sales volume.

## Reduced Risk

Lower chance of overstock and associated carrying costs.

## Customer Satisfaction

Improved availability of desired models enhances customer experience.



# Next Steps

1

## Integration

Integrate the model into dealer dashboards or API endpoints

2

## Retraining

Implement periodic retraining with new sales data

3

## Model Enhancement

Expand feature set and explore ensemble techniques to improve prediction accuracy.





# Thank You

## Questions?

All code and documentation are available in our GitHub repository:

*Scan to access the GitHub repo:*

[View GitHub Repository](#)



Made with **GAMMA**